



四川大學
SICHUAN UNIVERSITY

Database System Concepts

Advanced SQL

伍元凱

College of Computer Science (Software), Sichuan University

wuyk0@scu.edu.cn

2023/03/22

海納百川
有容乃大



SQL is a **domain-specific** language used for **managing** and **manipulating** relational databases. While SQL is a powerful tool for interacting with databases, it is not a **general-purpose** programming language.

- There exist queries that can be expressed in C, Java, or Python that cannot be expressed in SQL (regression analysis, hypothesis testing, time-series analysis, NLP, clustering, classification, and prediction)
- **Nondeclarative** action (printing, interacting, graphical user interface)

- **Dynamic SQL** refers to a programming technique that allows SQL statements to be constructed and executed at runtime, rather than being hard-coded into an application or stored procedure. Dynamic SQL is typically implemented using **string concatenation or substitution to build the SQL statement**, and then executing the resulting string as a single command.
- **Embedded SQL**, on the other hand, refers to the practice of embedding SQL statements directly into a host programming language, such as Java, C++, or COBOL. This approach allows SQL to be executed directly from within the host language, rather than having to call out to a separate SQL interpreter.

① Accessing SQL from a Programming Language

JDBC

Database Access from Python

OUBC

Embedded SQL and Database

② Functions and Procedures

③ Triggers

④ Recursive Queries

⑤ Advanced Aggregation Features

⑥ Ending


海纳百川 有容乃大

JDBC example

```
1 import java.sql.*;
2 public class ExampleJDBC {
3     public static void main(String[] args) {
4         try {
5             Class.forName("com.mysql.jdbc.Driver");
6             // Open a connection to the database
7             Connection conn = DriverManager.getConnection("jdbc:
                mysql://localhost:3306/mydatabase", "root", "
                mypassword");
8             // Create a statement object to execute SQL queries
9             Statement stmt = conn.createStatement();
10            // Execute a SELECT query and get the result set
11            ResultSet rs = stmt.executeQuery("SELECT * FROM mytable
                ");
12            // Iterate over the result set and print each row to
                the console
13            while (rs.next()) {
```

```
14         int id = rs.getInt("id");
15         String name = rs.getString("name");
16         double salary = rs.getDouble("salary");
17         System.out.println("id=" + id + ", name=" + name +
18                             ", salary=" + salary);
19     }
20     // Close the result set, statement, and connection
21     rs.close();
22     stmt.close();
23     conn.close();
24 } catch (SQLException e) {
25     e.printStackTrace();
26 } catch (ClassNotFoundException e) {
27     e.printStackTrace();
28 }
29 }
```

- ❶ register the JDBC driver `"Class.forName("com.mysql.jdbc.Driver")"`
- ❷ open a connection `"DriverManager.getConnection(url, username, password)"`
- ❸ create and execute statement `"stmt.executeQuery("SELECT * FROM mytable")"`
- ❹ iterate over the Result Set using while `"(rs.next())"`.
- ❺ close the ResultSet, Statement, and Connection objects.



Prepared Statement

```
1 PreparedStatement stmt = conn.prepareStatement("INSERT INTO  
    table_name (column1, column2, ...) VALUES (?, ?, ...)");  
2 stmt.setXXX(parameterIndex, value);  
3 stmt.setYYY(parameterIndex, value);  
4 ...  
5 stmt.executeUpdate();
```

- XXX, YYY, data type (String)
- parameterIndex, "?" in PreparedStatement
- value, the corresponding value of the parameterIndex

Real example

```
1 String sql = "INSERT INTO employees (name, salary, hire_date)  
    VALUES (?, ?, ?)";  
2 PreparedStatement stmt = conn.prepareStatement(sql);  
3 // Set the input parameters for the prepared statement  
4 stmt.setString(1, "John");  
5 stmt.setDouble(2, 50000.00);  
6 stmt.setDate(3, new java.sql.Date(System.currentTimeMillis()));
```


Dangerous query

```
1 String sourceAccount = request.getParameter("sourceAccount");
2 String destinationAccount = request.getParameter("
    destinationAccount");
3 double amount = Double.parseDouble(request.getParameter("amount"));
4
5 String updateQuery = "UPDATE accounts SET balance = balance - " +
    amount + " WHERE account_number = " + sourceAccount;
6 Statement stmt = conn.createStatement();
7 int numRowsAffected = stmt.executeUpdate(updateQuery);
8
9 updateQuery = "UPDATE accounts SET balance = balance + " + amount +
    " WHERE account_number = " + destinationAccount;
10 numRowsAffected = stmt.executeUpdate(updateQuery);
```

Inject "12345; DROP TABLE accounts; --" to "sourceAccount"

Injection Attack

```
1 UPDATE accounts SET balance = balance - 1000.0 WHERE account_number
    = 12345; DROP TABLE accounts; --;
```

Prevention by Prepared Statements

```
1 String sourceAccount = request.getParameter("sourceAccount");
2 String destinationAccount = request.getParameter("
    destinationAccount");
3 double amount = Double.parseDouble(request.getParameter("amount"));
4 String updateQuery = "UPDATE accounts SET balance = balance - ?
    WHERE account_number = ?";
5 PreparedStatement stmt = conn.prepareStatement(updateQuery);
6 stmt.setDouble(1, amount);
7 stmt.setString(2, sourceAccount);
8 int numRowsAffected = stmt.executeUpdate();
9 updateQuery = "UPDATE accounts SET balance = balance + ? WHERE
    account_number = ?";
10 stmt = conn.prepareStatement(updateQuery);
11 stmt.setDouble(1, amount);
12 stmt.setString(2, destinationAccount);
13 numRowsAffected = stmt.executeUpdate();
```

If an attacker tries to enter the same input as before (12345; DROP TABLE accounts; –) as the source account number, the prepared statement will treat it as a string value and escape any special characters

- Prepared statements allow the DBMS to **cache the execution plan** and reuse it across multiple executions of the same statement.
- In contrast, when a query is executed without using a prepared statement, the DBMS must parse and optimize the query every time it is executed.



Callable Statements

```
1 String procedure = "{CALL my_stored_procedure(?, ?, ?)}";
2 CallableStatement stmt = conn.prepareCall(procedure);
3
4 // set input parameters
5 stmt.setString(1, "John");
6 stmt.setString(2, "Doe");
7 stmt.setInt(3, 30);
8
9 // register output parameter
10 stmt.registerOutParameter(3, Types.INTEGER);
11
12 // execute the stored procedure
13 stmt.execute();
14
15 // retrieve the output parameter
16 int result = stmt.getInt(3);
```

Invocation of SQL stored procedures and functions.

database product name, version, and driver version:

```
1 DatabaseMetaData metadata = connection.getMetaData();
2 String productName = metadata.getDatabaseProductName();
3 String productVersion = metadata.getDatabaseProductVersion();
4 String driverVersion = metadata.getDriverVersion();
```

Get a list of tables in the database:

```
1 DatabaseMetaData metadata = connection.getMetaData();
2 ResultSet tables = metadata.getTables(null, null, null, new String
    [] {"TABLE"});
3 while (tables.next()) {
4     String tableName = tables.getString("TABLE_NAME");
5     // process the table name
6 }
```

Get information about the columns of a table:

```
1 DatabaseMetaData metadata = connection.getMetaData();
2 ResultSet columns = metadata.getColumns(null, null, "my_table",
    null);
3 while (columns.next()) {
4     String columnName = columns.getString("COLUMN_NAME");
5     String dataType = columns.getString("DATA_TYPE");
6     int columnSize = columns.getInt("COLUMN_SIZE");
7     // process the column information
8 }
```

Get information about the primary keys of a table:

```
1 DatabaseMetaData metadata = connection.getMetaData();
2 ResultSet primaryKeys = metadata.getPrimaryKeys(null, null, "
    my_table");
3 while (primaryKeys.next()) {
4     String columnName = primaryKeys.getString("COLUMN_NAME");
5     int keySeq = primaryKeys.getInt("KEY_SEQ");
6     // process the primary key information
7 }
```

Dealing Transactions:

```
1 // assume conn is a valid Connection object
2 try {
3     // set auto-commit to false to start a transaction
4     conn.setAutoCommit(false);
5     // execute some SQL statements within the transaction
6     Statement stmt = conn.createStatement();
7     stmt.executeUpdate("UPDATE my_table SET my_column = 'value'
8                        WHERE id = 1");
9     stmt.executeUpdate("DELETE FROM my_table WHERE id = 2");
10    // commit the transaction if all statements executed
11    // successfully
12    conn.commit();
13 } catch (SQLException e) {
14     // rollback the transaction if any statement fails
15     conn.rollback();
16     e.printStackTrace();
17 } finally {
18
19
20 // restore auto-commit to true after the transaction completes
21     conn.setAutoCommit(true);
22 }
```

We first set the **autoCommit** property of the Connection object to false to start a transaction. If all statements execute successfully, we call the **commit** method to commit the transaction. If any statement fails, we catch the **SQLException** and call the **rollback** method to roll back the transaction.



..oo | 目录

① Accessing SQL from a Programming Language

JDBC

Database Access from Python

OUBC

Embedded SQL and Database

② Functions and Procedures

③ Triggers

④ Recursive Queries

⑤ Advanced Aggregation Features

⑥ Ending

海納百川 有容乃大

```
1 import psycopg2
2 # establish a database connection
3 conn = psycopg2.connect(
4     host="localhost",
5     database="my_database",
6     user="my_username",
7     password="my_password"
8 )
9 # create a cursor object
10 cur = conn.cursor()
11 # prepare a SQL statement
12 query = "SELECT name, age FROM my_table WHERE id = %s"
13 cur.execute("PREPARE prepared_query AS {}".format(query))
14 # execute the prepared statement with different parameters
15 id1 = 1
16 cur.execute("EXECUTE prepared_query (%s)", (id1,))
17 print(cur.fetchone())

18 cur.execute("DEALLOCATE prepared_query")
19 # commit the transaction and close the cursor and connection
20 conn.commit()
21 cur.close()
```

- `psycopg2.connect`: Establish a database connection
- `cursor`: Execute SQL queries and fetch the results.
- `execute`: passing in the query string and any necessary parameters.
- `commit`: commit the transaction.



① Accessing SQL from a Programming Language

JDBC

Database Access from Python

OUBC

Embedded SQL and Database

② Functions and Procedures

③ Triggers

④ Recursive Queries

⑤ Advanced Aggregation Features

⑥ Ending

海纳百川 有容乃大

```
1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <sql.h>
4 #include <sqlext.h>
5 #define BUFFER_SIZE 256
6 int main() {
7     SQLHENV env;
8     SQLHDBC dbc;
9     SQLHSTMT stmt;
10    SQLCHAR query[BUFFER_SIZE];
11    SQLRETURN ret;
12    SQLINTEGER id;
13    SQLCHAR name[BUFFER_SIZE];
14    SQLINTEGER age;
15    // Allocate environment handle
16    SQLAllocHandle(SQL_HANDLE_ENV, SQL_NULL_HANDLE, &env);
17    // Set ODBC version to 3
```

```
18  SQLSetEnvAttr(env, SQL_ATTR_ODBC_VERSION, (SQLPOINTER)
    SQL_OV_ODBC3, 0);
19  // Allocate connection handle
20  SQLAllocHandle(SQL_HANDLE_DBC, env, &dbc);
21  // Connect to database
22  SQLConnect(dbc, (SQLCHAR*) "DSN=mydatabase", SQL_NTS, (SQLCHAR
    *) "", SQL_NTS, (SQLCHAR*) "", SQL_NTS);
23  // Allocate statement handle
24  SQLAllocHandle(SQL_HANDLE_STMT, dbc, &stmt);
25  // Prepare SQL query
26  SQLPrepare(stmt, (SQLCHAR*) "SELECT * FROM mytable WHERE id = ?
    ", SQL_NTS);
27  // Bind input parameter to statement
28  SQLBindParameter(stmt, 1, SQL_PARAM_INPUT, SQL_C_LONG,
    SQL_INTEGER, 0, 0, &id, 0, NULL);
29  id = 1;
```

```
1 // Execute statement
2 SQLExecute(stmt);
3 // Bind result columns to variables
4 SQLBindCol(stmt, 1, SQL_C_CHAR, name, BUFFER_SIZE, NULL);
5 SQLBindCol(stmt, 2, SQL_C_LONG, &age, 0, NULL);
6 // Fetch results
7 while (SQLFetch(stmt) == SQL_SUCCESS) {
8     printf("Name: %s\n", name);
9     printf("Age: %d\n", age);
10 }
11 // Clean up resources
12 SQLFreeHandle(SQL_HANDLE_STMT, stmt);
13 SQLDisconnect dbc;
14 SQLFreeHandle(SQL_HANDLE_DBC, dbc);
15 SQLFreeHandle(SQL_HANDLE_ENV, env);
16 return 0;
17 }
```


目录

① Accessing SQL from a Programming Language

JDBC

Database Access from Python

OUBC

Embedded SQL and Database

② Functions and Procedures

③ Triggers

④ Recursive Queries

⑤ Advanced Aggregation Features

⑥ Ending

海纳百川 有容乃大

- An embedded database is a database management system that is **integrated into an application, rather than running as a separate server or service**. This means that the database engine and the application are packaged together, and the application communicates directly with the database engine to perform database operations.
- Embedded SQL is a technique for integrating SQL statements into host language code. With embedded SQL, developers can **write SQL statements directly in their host language code**, and use the language's native interface to execute those statements against a database.

① Accessing SQL from a Programming Language

② Functions and Procedures

Declaring and Invoking SQL
Functions and Procedures
Language Constructs for
Procedures and Functions
External Language Routines

③ Triggers

④ Recursive Queries

⑤ Advanced Aggregation Features

⑥ Ending

海纳百川 有容乃大

A complex function

```
1 CREATE FUNCTION GetTopSellingProductsByCategory(category VARCHAR
    (50),
2 limit INTEGER)
3 RETURNS TABLE(ProductName VARCHAR(50), SalesAmount DECIMAL
    (10,2))
4 RETURNS TABLE
5     (SELECT p.ProductName, SUM(od.Quantity * od.UnitPrice) AS
        SalesAmount
6     FROM Products p
7     INNER JOIN Categories c ON p.CategoryID = c.CategoryID
8     INNER JOIN OrderDetails od ON p.ProductID = od.ProductID
9     INNER JOIN Orders o ON od.OrderID = o.OrderID
10    WHERE c.CategoryName = $1
11    GROUP BY p.ProductName
12    ORDER BY SalesAmount DESC);
13
14 SELECT * FROM GetTopSellingProductsByCategory('Beverages', 3);
```

A complex procedure

```
1 CREATE PROCEDURE calculate_sales_revenue(  
2     IN category_name VARCHAR(50),  
3     IN start_date DATE,  
4     IN end_date DATE,  
5     OUT total_revenue DECIMAL(10,2)  
6 )  
7 BEGIN  
8     SELECT SUM(p.price * oi.quantity) INTO total_revenue  
9     FROM products p  
10    JOIN order_items oi ON p.product_id = oi.product_id  
11    JOIN orders o ON oi.order_id = o.order_id  
12    JOIN categories c ON p.category_id = c.category_id  
13    WHERE c.category_name = category_name  
14    AND o.order_date BETWEEN start_date AND end_date;  
15 END  
16  
17 CALL calculate_sales_revenue('Electronics', '2022-01-01', '  
  
2022-12-31', @total_revenue);
```

6 Ending

```
1 CREATE PROCEDURE calculate_factorial(  
2     IN number INT,  
3     OUT factorial INT  
4 )  
5 BEGIN  
6     SET factorial = 1;  
7     SET @counter = 1;  
8     WHILE (@counter <= number) DO  
9         SET factorial = factorial * @counter;  
10        SET @counter = @counter + 1;  
11    END WHILE;  
12 END  
13  
14 SET @result = 0;  
15 CALL calculate_factorial(5, @result);  
16 SELECT @result;
```

```
1 CREATE PROCEDURE find_min_max_salary(  
2     OUT min_salary DECIMAL(10, 2),  
3     OUT max_salary DECIMAL(10, 2)  
4 )  
5 BEGIN  
6     DECLARE done INT DEFAULT FALSE;  
7     DECLARE cur_salary DECIMAL(10, 2);  
8     DECLARE cur CURSOR FOR  
9         SELECT salary FROM employees;  
10    DECLARE CONTINUE HANDLER FOR NOT FOUND SET done = TRUE;  
11    OPEN cur;  
12    SET min_salary = 99999999.99;  
13    SET max_salary = 0.00;  
14    REPEAT  
15        FETCH cur INTO cur_salary;  
16        IF NOT done THEN  
17            IF cur_salary < min_salary THEN
```



```
18         SET min_salary = cur_salary;
19     END IF;
20     IF cur_salary > max_salary THEN
21         SET max_salary = cur_salary;
22     END IF;
23 END IF;
24 UNTIL done END REPEAT;
25 CLOSE cur;
26 END
27
28 CALL find_min_max_salary(@min_salary, @max_salary);
29 SELECT @min_salary, @max_salary;
```

- **DECLARE**: This keyword is used to declare local variables and cursors within the procedure.
- **CURSOR**: A cursor is a mechanism that allows you to fetch rows from a result set one at a time.
- **HANDLER**: A handler is a type of construct that allows you to define an action to take when a certain condition occurs.
- **OPEN**: This keyword is used to open the cursor and prepare it for fetching rows.
- **FETCH**: This keyword is used to fetch the next row from the cursor and assign its value to a local variable.
- **REPEAT**: This keyword is used to define the start of a loop that will repeat until a certain condition is met.
- **UNTIL**: This keyword is used to specify the condition that must be met in order for the loop to exit.
- **END REPEAT**: This keyword is used to define the end of the loop.
- **CLOSE**: This keyword is used to close the cursor once we're done fetching rows from it.

```
1 CREATE PROCEDURE update_product_prices(  
2     IN price_change DECIMAL(10,2)  
3 )  
4 BEGIN  
5     DECLARE product_id INT;  
6  
7     FOR product_id IN (SELECT id FROM products) DO  
8         UPDATE products SET price = price + price_change WHERE id =  
9             product_id;  
10    END FOR;  
11  
12 CALL update_product_prices(10.00);
```

The procedure will then update the prices of all products in the database by adding 10.00 to each product's price.

① Accessing SQL from a Programming Language

② Functions and Procedures

Declaring and Invoking SQL
Functions and Procedures
Language Constructs for
Procedures and Functions
External Language Routines

③ Triggers

④ Recursive Queries

⑤ Advanced Aggregation Features

⑥ Ending

海纳百川 有容乃大

```
1 CREATE [OR REPLACE] FUNCTION function_name ([argument_list])
2     RETURNS return_type
3     LANGUAGE language_name
4     [EXTERNAL NAME 'external_name']
5 CREATE [OR REPLACE] PROCEDURE procedure_name (in [argument_list],
6     out [argument_list])
7     LANGUAGE language_name
8     [EXTERNAL NAME 'external_name']
```

- **CREATE [OR REPLACE] FUNCTION|PROCEDURE** is used to create a new function or procedure or replace an existing one.
- **function_name, procedure_name** is the name of the function or procedure.
- **LANGUAGE language_name** specifies the programming language used for the function or procedure.
- **EXTERNAL NAME 'external_name'** is used to specify the name of the external function or procedure that is being called by the SQL function or procedure.

```
1 CREATE FUNCTION compute_average_age (IN table_name TEXT, OUT
   average_age INTEGER)
2 RETURNS RECORD
3 LANGUAGE 'plpythonu'
4 AS $$
5     import plpy
6     conn = plpy.connect()
7     query = f"SELECT age FROM {table_name}"
8     age_result = conn.execute(query)
9     # Calculate the average age
10    total_age = 0
11    num_rows = 0
12    for row in age_result:
13        total_age += row['age']
14        num_rows += 1
15    average_age = total_age // num_rows
16    # Return the result as a record
17
18    return {'average_age': average_age}
19    $$;
```

Triggers are used in SQL to automatically execute a set of actions when a specified event (insertion, update and delete) occurs on a table or view. Triggers are useful because they allow you to automate complex data management tasks that would be **difficult or time-consuming to perform manually**. They can also help you **maintain data integrity**.

```
1 CREATE TRIGGER update_customer_stats
2 AFTER INSERT ON orders
3 FOR EACH ROW
4 EXECUTE PROCEDURE update_customer_total();
5
6 CREATE PROCEDURE update_customer_total() RETURNS trigger AS $$
7 BEGIN
8     UPDATE customer_stats
9     SET total_order_amount = total_order_amount + NEW.order_amount
10    WHERE customer_id = NEW.customer_id;
11    RETURN NEW;
12 END;
13 $$ LANGUAGE plpgsql;
```

When Not to Use Triggers

- ① Accessing SQL from a Programming Language
- ② Functions and Procedures
- ③ Triggers
Triggers in SQL
- ④ Recursive Queries
- ⑤ Advanced Aggregation Features
- ⑥ Ending

海纳百川 有容乃大

Standard Trigger

```
1 CREATE [OR REPLACE] TRIGGER trigger_name
2 {BEFORE | AFTER} {INSERT | UPDATE | DELETE}
3 ON table_name
4 [REFERENCING OLD AS old NEW AS new]
5 [FOR EACH ROW]
6 [WHEN (condition)]
7 trigger_body;
```

- **CREATE [OR REPLACE] TRIGGER** - This creates a new trigger or replaces an existing one with the same name.
- **BEFORE | AFTER** - This specifies whether the trigger should fire before or after the specified event occurs on the table.
- **[REFERENCING]** - This clause is used to define aliases for the OLD and NEW variables that are used in the trigger action.
- **[WHEN (condition)]** - This clause is used to specify a condition that must be true for the trigger to fire.

Table 1: Employees table

id	name	salary	departments
1	John Doe	50000	Sales
2	Jane Doe	60000	HR
3	Bob Smith	70000	HR

```
1 CREATE TRIGGER prevent_salary_updates
2 AFTER UPDATE ON employees
3 FOR EACH ROW
4 BEGIN
5     -- Check if the employee is in the Sales department
6     IF NEW.department = 'Sales' THEN
7         -- If they are, prevent the update to the salary column
8         SET NEW.salary = OLD.salary;
9     END IF;
10 END;
11
12 UPDATE employees SET salary = 55000 WHERE id = 1;
```

```
1 SELECT * FROM employees;
```

Table 2: Employees table

id	name	salary	departments
1	John Doe	50000	Sales
2	Jane Doe	60000	HR
3	Bob Smith	70000	HR

Using **before** to prevent errors

```
1 CREATE TRIGGER enforce_positive_total_price_and_discount
2 BEFORE INSERT ON orders
3 FOR EACH ROW
4 BEGIN
5     IF NEW.total_price <= 0 THEN
6         SIGNAL SQLSTATE '45000'
7         SET MESSAGE_TEXT = 'Total price must be a positive number';
8     END IF;
9     IF NEW.discount < 0 THEN
10        SIGNAL SQLSTATE '45000'
11        SET MESSAGE_TEXT = 'Discount must be a non-negative number';
12    END IF;
13 END;
14
15 INSERT INTO orders (id, total_price, discount) VALUES (1, -100, 20)
16      ;
17 ERROR 1644 (45000): Total price must be a positive number
```

- `alter trigger ... disable/enable`: disable and enable a specific trigger.
- `drop trigger ...`: removes a trigger permanently.



① Accessing SQL from a Programming Language

③ Triggers

Triggers in SQL

④ Recursive Queries

5 Advanced Aggregation Features

6 Ending

Cascade Deletion

```
1 CREATE TABLE departments (  
2   id INT PRIMARY KEY,  
3   name VARCHAR(50)  
4 );  
5  
6 CREATE TABLE employees (  
7   id INT PRIMARY KEY,  
8   name VARCHAR(50),  
9   department_id INT,  
10  FOREIGN KEY (department_id) REFERENCES departments(id)  
11 );  
12  
13 CREATE TRIGGER delete_employees  
14 ON departments  
15 FOR DELETE  
16 AS  
17 BEGIN  
  
18   DELETE FROM employees  
19   WHERE department_id IN (SELECT id FROM DELETED);  
20 END;
```

Maintain materialized views

```
1 CREATE TABLE sales (  
2   id INT PRIMARY KEY,  
3   customer_id INT,  
4   amount DECIMAL(10, 2),  
5   date DATE  
6 );  
7  
8 CREATE MATERIALIZED VIEW customer_sales AS  
9 SELECT customer_id, SUM(amount) AS total_sales  
10 FROM sales  
11 GROUP BY customer_id;  
12  
13 CREATE TRIGGER update_customer_sales  
14 AFTER INSERT OR UPDATE OR DELETE ON sales  
15 FOR EACH STATEMENT  
16 AS  
17 BEGIN  
  
18   REFRESH MATERIALIZED VIEW customer_sales;  
19 END;
```


-

A poor designed trigger with performance issue

```
1 CREATE TABLE customers (  
2   id INTEGER PRIMARY KEY,  
3   name VARCHAR(255),  
4   email VARCHAR(255),  
5   total_spent DECIMAL(10,2)  
6 );  
7  
8 CREATE TABLE orders (  
9   id INTEGER PRIMARY KEY,  
10  customer_id INTEGER,  
11  amount DECIMAL(10,2),  
12  order_date DATE  
13 );  
14  
15 CREATE TRIGGER update_customer_totals  
16 AFTER INSERT ON orders  
17 FOR EACH ROW
```

```
18 BEGIN
19  UPDATE customers
20  SET total_spent = (SELECT SUM(amount) FROM orders WHERE
                     customer_id = NEW.customer_id)
21  WHERE id = NEW.customer_id;
22 END;
```

A modified one

```
1 CREATE TRIGGER update_customer_totals
2 AFTER INSERT ON orders
3 FOR EACH ROW
4 BEGIN
5     UPDATE customers
6     SET total_spent = total_spent + NEW.amount
7     WHERE id = NEW.customer_id;
8 END;
```

This trigger simply **adds the new order's amount** to the customer's existing total spent, rather than **recalculating the total spent** for every customer every time a new order is inserted.

Employees table

employee_id	employee_name	supervisor_id
1	Alice	NULL
2	Bob	1
3	Charlie	2
4	Dave	2
5	Eve	1
6	Frank	5

In this example, Alice is the **CEO of the organization and has no supervisor**. Bob and Eve report directly to Alice, while Charlie and Dave report to Bob, and Frank reports to Eve.

Resulted table

level	name
1	Alice
2	-Bob
3	-Charlie
3	-Dave
2	-Eve
3	-Frank



.. | 目录

① Accessing SQL from a
Programming Language

② Functions and Procedures

③ Triggers

④ Recursive Queries

Transitive Closure Using
Iteration
Recursion in SQL

⑤ Advanced Aggregation
Features

⑥ Ending

海纳百川 有容乃大

Without recursive query

```
1 WITH
2   level_one AS (
3     SELECT employee_name, 1 AS level
4     FROM employee
5     WHERE supervisor_id IS NULL
6   ),
7   level_two AS (
8     SELECT CONCAT('-', employee_name) AS employee_name, 2 AS level
9     FROM employee
10    WHERE supervisor_id IN (SELECT employee_id FROM level_one)
11  ),
12  level_three AS (
13    SELECT CONCAT('--', employee_name) AS employee_name, 3 AS level
14    FROM employee
15    WHERE supervisor_id IN (SELECT employee_id FROM level_two)
16  )
17 SELECT employee_name, level
```



```
18 FROM level_one
19 UNION ALL
20 SELECT employee_name, level
21 FROM level_two
22 UNION ALL
23 SELECT employee_name, level
24 FROM level_three;
```

- ① The first CTE, `level_one`, extracts the employees who do not have a supervisor.
- ② The second and third CTE, `level_two` and `level_three`, extracts the employees who have a supervisor in `level_one` and `level_two`.
- ③ Finally, we combine the results of all three CTEs using **UNION ALL** to produce the desired output, which includes the employee names and their respective levels.



Recursion in SQL | 目录

① Accessing SQL from a
Programming Language

② Functions and Procedures

③ Triggers

④ Recursive Queries

Transitive Closure Using
Iteration
Recursion in SQL

⑤ Advanced Aggregation
Features

⑥ Ending

海納百川 有容乃大

Recursive query

```
1 WITH RECURSIVE org_chart AS (  
2   SELECT employee_id, employee_name, supervisor_id, 1 AS level  
3   FROM employees  
4   WHERE supervisor_id IS NULL  
5   UNION ALL  
6   SELECT e.employee_id, e.employee_name, e.supervisor_id, oc.level  
       + 1  
7   FROM employees e  
8   JOIN org_chart oc ON e.supervisor_id = oc.employee_id  
9 )  
10 SELECT CONCAT(REPEAT('-', level - 1), employee_name) AS name, level  
11 FROM org_chart  
12 ORDER BY level, name;
```


General Form

```
1 WITH RECURSIVE cte_name AS (  
2     -- Anchor member  
3     SELECT ...  
4     UNION ALL  
5     -- Recursive member  
6     SELECT ...  
7     FROM cte_name  
8     WHERE ...  
9 )  
10 SELECT ...  
11 FROM cte_name  
12 WHERE ...
```

In summary, a recursive SQL query uses a CTE to define a set of rules for generating a result set that includes **multiple levels** of data. The query includes an **anchor member that defines the starting point for the recursion**, and a **recursive member** that generates new rows based on the results of previous iterations.

- **WITH RECURSIVE:** This clause starts the CTE and tells the database that it should perform a recursive query.
- **cte_name:** This is the name you give to the CTE. You'll use this name later in the query to reference the results of the CTE.
- **UNION ALL:** This clause combines the anchor member and the recursive member of the CTE.

- Aggregate functions: Recursive queries cannot use aggregate functions such as SUM, AVG, COUNT, MAX, MIN, etc.
- GROUP BY: Similarly, recursive queries cannot use the GROUP BY clause in the recursive part of the query.
- Non-deterministic functions: Recursive queries cannot use non-deterministic functions such as RAND(), NEWID(), or GETDATE() in the recursive part of the query.
- Set difference whose right-hand side uses the recursive view.



- ① Accessing SQL from a Programming Language
- ② Functions and Procedures
- ③ Triggers
- ④ Recursive Queries

⑤ Advanced Aggregation Features

- Ranking
- Windowing
- Pivoting
- Rollup and Cube

⑥ Ending

海纳百川 有容乃大

Ranking is a common technique used in SQL to **assign a rank** or order to rows within a result set, based on one or more criteria.

General Form

```
1 SELECT col1, col2, ..., RANK() OVER (ORDER BY col3 DESC) AS rank  
2 FROM table_name;
```

- **col1, col2, etc.** represent the columns you want to select from the table.
- **RANK() OVER** is the ranking function. You can replace **RANK()** with other ranking functions such as **DENSE_RANK()** or **ROW_NUMBER()** if needed.
- **ORDER BY col3 DESC** specifies the column(s) to use for ranking. In this case, we're using col3 in descending order .
- **AS rank** is an optional alias for the ranking column.

Sales

ID	Sales	Region	Date
1	1000	West	2022-01-01
2	750	East	2022-01-02
3	1200	South	2022-01-03
4	NULL	West	2022-01-04
5	1500	North	2022-01-05
6	800	East	2022-01-06
7	1100	South	2022-01-07
8	NULL	North	2022-01-08
9	1300	West	2022-01-09
10	900	East	2022-01-10

```
1 SELECT ID, Sales, RANK() OVER (ORDER BY Sales DESC NULLS LAST) AS  
   SalesRank  
2 FROM Sales;
```

Table 3: Sales Rank Table

ID	Sales	SalesRank
5	1500	1
3	1200	2
9	1300	3
7	1100	4
1	1000	5
6	800	6
10	900	7
2	750	8
4	NULL	8
8	NULL	8



Without ranking

```
1 SELECT
2   ID,
3   Sales,
4   (
5     SELECT COUNT(*)
6     FROM
7       (SELECT DISTINCT Sales FROM SalesTable WHERE Sales IS NOT
8         NULL) AS temp
9     WHERE
10       SalesTable.Sales < temp.Sales OR
11       (SalesTable.Sales IS NULL AND temp.Sales IS NOT NULL)
12   ) + 1 AS SalesRank
13 FROM
14   SalesTable
15 ORDER BY
16   SalesRank ASC,
17   ID ASC;
```

```
1 SELECT ID, Sales, Region, Date,  
2     RANK() OVER (PARTITION BY Region ORDER BY Sales DESC NULLS  
3                 LAST) AS Region_Rank  
4 FROM sales;
```

Equal query

```
1 SELECT s1.ID, s1.Sales, s1.Region, s1.Date,  
2     COUNT(s2.ID) + 1 AS Region_Rank  
3 FROM sales s1  
4 LEFT JOIN sales s2  
5     ON s1.Region = s2.Region AND s1.Sales < s2.Sales  
6 GROUP BY s1.ID, s1.Sales, s1.Region, s1.Date  
7 ORDER BY s1.Region, s1.Sales DESC;
```

Sales Table with Region Ranks

ID	Sales	Region	Date	Region_Rank
9	1300	West	2022-01-09	1
1	1000	West	2022-01-01	2
4	NULL	West	2022-01-04	2
10	900	East	2022-01-10	1
6	800	East	2022-01-06	2
2	750	East	2022-01-02	3
5	1500	North	2022-01-05	1
8	NULL	North	2022-01-08	1
3	1200	South	2022-01-03	1
7	1100	South	2022-01-07	2

RANK() in subquery

```
1 SELECT ID, Sales, Region, Date
2 FROM (
3     SELECT ID, Sales, Region, Date, RANK() OVER (ORDER BY Sales DESC)
4         AS Sales_Rank
5     FROM Sales
6 )
7 WHERE Sales_Rank = 1;
```

Sales Table with Region Ranks

ID	Sales	Region	Date	Sales_Rank
5	1500	North	2022-01-05	1

Some other functions that can be used in place of `RANK()` in SQL:

- **DENSE_RANK()**: This function is similar to `RANK()`, but it does not skip any ranks if there are ties. Instead, it assigns the same rank to all tied rows and then continues assigning ranks in sequential order.
- **ROW_NUMBER()**: This function simply assigns a unique number to each row in the result set, starting from 1.
- **NTILE()**: This function divides the result set into a specified number of "tiles" or groups.
- **PERCENT_RANK()**: This function returns the relative rank of each row within the result set as a percentage.
- **CUME_DIST()**: This function returns the cumulative distribution of each row within the result set.
- ...


```
1 SELECT *, NTILE(4) OVER(ORDER BY Sales) as Quartile
2 FROM Sales;
```

ID	Sales	Region	Date	Quartile
2	750	East	2022-01-02	1
6	800	East	2022-01-06	1
10	900	East	2022-01-10	1
1	1000	West	2022-01-01	2
7	1100	South	2022-01-07	2
3	1200	South	2022-01-03	2
9	1300	West	2022-01-09	3
5	1500	North	2022-01-05	3
4	NULL	West	2022-01-04	4
8	NULL	North	2022-01-08	4

Sales by Region, Date, and Quartile

4 Recursive Queries

Rollup and Cube

6 Ending



suppose you want to calculate the 30-day moving average of a stock's closing price. You could use the following SQL query:

```
1 SELECT date, close_price, AVG(close_price) OVER (  
2     ORDER BY date  
3     ROWS BETWEEN 29 PRECEDING AND CURRENT ROW  
4 ) AS moving_avg  
5 FROM stock_prices
```

```
1 <window function>([expression]) OVER (  
2     [PARTITION BY partition_expression, ... ]  
3     [ORDER BY sort_expression [ASC | DESC], ... ]  
4     [ROWS <frame specification>]  
5 )
```

- **window function**: SUM, AVG, MAX, MIN, etc.
- **expression**: attribute/column
- **frame specification**: defines the subset of rows to include in the window function calculation. There are three types of frame specification:

date	region	product	sales_amount
2022-01-01	North	Product A	1000.00
2022-01-01	North	Product B	1500.00
2022-01-02	North	Product A	2000.00
2022-01-02	North	Product B	2500.00
2022-01-01	South	Product A	3000.00
2022-01-01	South	Product B	3500.00
2022-01-02	South	Product A	4000.00
2022-01-02	South	Product B	4500.00

Rolling Average

```
1 SELECT date, region, product, sales_amount, AVG(sales_amount) OVER(  
    PARTITION BY region, product ORDER BY date ROWS BETWEEN 1  
    PRECEDING AND CURRENT ROW) as rolling_avg  
2 FROM table_name;  
3 )
```

date	region	product	sales_amount	rolling_avg
2022-01-01	North	Product A	1000.00	1000.00
2022-01-01	North	Product B	1500.00	1500.00
2022-01-02	North	Product A	2000.00	1500.00
2022-01-02	North	Product B	2500.00	2000.00
2022-01-01	South	Product A	3000.00	3000.00
2022-01-01	South	Product B	3500.00	3500.00
2022-01-02	South	Product A	4000.00	3500.00
2022-01-02	South	Product B	4500.00	4000.00

Rank the sales_amount

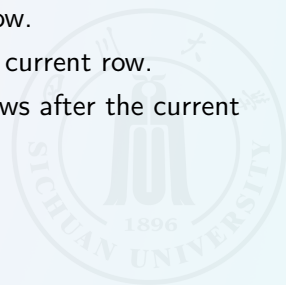
```
1 SELECT date, region, product, sales_amount, RANK() OVER(PARTITION
    BY region, product ORDER BY sales_amount DESC) as rank
2 FROM table_name;
```

date	region	product	sales_amount	rank
2022-01-01	North	Product A	1000.00	2
2022-01-01	North	Product B	1500.00	2
2022-01-02	North	Product A	2000.00	1
2022-01-02	North	Product B	2500.00	1
2022-01-01	South	Product A	3000.00	2
2022-01-01	South	Product B	3500.00	2
2022-01-02	South	Product A	4000.00	1
2022-01-02	South	Product B	4500.00	1


Frame specification:

ROWS BETWEEN: Specifies a range of rows relative to the current row to include in the window frame. The possible values are:

- **UNBOUNDED PRECEDING:** Includes all rows before the current row.
- **n PRECEDING:** Includes the n rows before the current row.
- **CURRENT ROW:** Includes only the current row.
- **n FOLLOWING:** Includes the n rows after the current row.
- **UNBOUNDED FOLLOWING:** Includes all rows after the current row.



- **RANGE BETWEEN:** Specifies a range of values relative to the current row to include in the window frame. **This is only applicable for window functions that use an ordering clause. The possible values are the same as ROWS BETWEEN.**
- **GROUPS BETWEEN:** Specifies a range of peer groups relative to the current row to include in the window frame.
- **UNBOUNDED PRECEDING:** Includes all rows before the current row.
- **UNBOUNDED FOLLOWING:** Includes all rows after the current row.



① Accessing SQL from a
Programming Language

② Functions and Procedures

③ Triggers

④ Recursive Queries

⑤ Advanced Aggregation
Features

Ranking

Windowing

Pivoting

Rollup and Cube

⑥ Ending

海纳百川 有容乃大

Pivoting in SQL can be very useful in situations where you need to transform data from **rows (values)** into **columns (attributes)** to make it easier to analyze.

Pivot in SQL

1. Use `GROUP BY` to combine rows
2. Use `FILTER` to pick rows per column

Year	Month	Revenue
2016	1	1
2016	2	23
2016	3	345
2016
2016	12	1234

Year	Jan	Feb	Mar	...	Dec
2016	1	23	345	...	1234

`SUM(...)` `FILTER(WHERE ...)`

`SUM(revenue)` `FILTER(WHERE month=1)`

`SUM(revenue)` `FILTER(WHERE month=2)`

`SUM(revenue)` `FILTER(WHERE month=3)`


`SUM(revenue)` `FILTER(WHERE month=...)`

`SUM(revenue)` `FILTER(WHERE month=12)`



Month	Product	Sales	Expenses
January	Product A	1000	500
January	Product B	1500	750
February	Product A	1200	600
February	Product B	1800	900
March	Product A	800	400
March	Product B	1200	600

Sales and expenses by product and month



Without Pivoting

```
1 SELECT
2     Month,
3     SUM(CASE WHEN Product = 'Product A' THEN Sales ELSE 0 END) AS '
        Product A Sales',
4     SUM(CASE WHEN Product = 'Product A' THEN Expenses ELSE 0 END)
        AS 'Product A Expenses',
5     SUM(CASE WHEN Product = 'Product B' THEN Sales ELSE 0 END) AS '
        Product B Sales',
6     SUM(CASE WHEN Product = 'Product B' THEN Expenses ELSE 0 END)
        AS 'Product B Expenses'
7 FROM
8     SalesExpenses
9 GROUP BY
10    Month;
```

Pivoting

```
1 SELECT Month,
2     [Product A Sales], [Product A Expenses],
3     [Product B Sales], [Product B Expenses]
4 FROM (
5     SELECT Month, Product, Sales, Expenses
6     FROM table1
7 ) AS t
8 PIVOT (
9     SUM(Sales) FOR Product IN ([Product A], [Product B])
10 ) AS SalesPivot
11 PIVOT (
12     SUM(Expenses) FOR Product IN ([Product A], [Product B])
13 ) AS ExpensesPivot
```

Month	Product A Sales	Product A Expenses	Product B Sales	Product B Expenses
January	1000	500	1500	750
February	1200	600	1800	900
March	800	400	1200	600

General Form

```
1 SELECT <non-pivoted column(s)>
2     , [first pivoted column] AS <column name>
3     , [second pivoted column] AS <column name>
4     , ...
5 FROM
6     (<source table>)
7 PIVOT
8 (
9     <aggregation function>(<value column>)
10    FOR
11    [<column to pivot on>]
12    IN ( [first pivoted column], [second pivoted column], ... )
13 ) AS <alias for the pivot table>
```

① Accessing SQL from a
Programming Language

② Functions and Procedures

③ Triggers

④ Recursive Queries

⑤ Advanced Aggregation
Features

Ranking

Windowing

Pivoting

Rollup and Cube

⑥ Ending

海纳百川 有容乃大

ROLLUP is a SQL keyword used to generate **subtotals** and grand totals of data. It generates **multiple grouping sets** by specifying columns on which data should be aggregated.

Region	Country	Product	SalesAmount
North	USA	Laptop	2000
North	USA	Phone	1500
North	Canada	Laptop	1800
North	Canada	Phone	1200
South	Brazil	Laptop	2500
South	Brazil	Phone	3000
South	Argentina	Laptop	1800
South	Argentina	Phone	2200

Sales

```
1 SELECT Region, Country, Product, SUM(SalesAmount) as TotalSales
2 FROM Sales
3 GROUP BY ROLLUP (Region, Country, Product)
```

Region	Country	Product	TotalSales
North	Canada	Laptop	1,800
North	Canada	Phone	1,200
North	USA	Laptop	2,000
North	USA	Phone	1,500
North	NULL	NULL	6,500
South	Argentina	Laptop	1,800
South	Argentina	Phone	2,200
South	Brazil	Laptop	2,500
South	Brazil	Phone	3,000
South	NULL	NULL	9,500
NULL	NULL	NULL	16,000

Sales by Region, Country, and Product

Without Rollup

```
1 SELECT Region, Country, Product, SUM(SalesAmount) as TotalSales
2 FROM Sales
3 GROUP BY Region, Country, Product
4
5 UNION ALL
6
7 SELECT Region, Country, NULL as Product, SUM(SalesAmount) as
   TotalSales
8 FROM Sales
9 GROUP BY Region, Country
10
11 UNION ALL
12
13 SELECT Region, NULL as Country, NULL as Product, SUM(SalesAmount)
   as TotalSales
14 FROM Sales
15 GROUP BY Region
```

Multiple Rollup

```
1 SELECT Region, Country, SUM(SalesAmount) AS TotalSales
2 FROM Sales
3 GROUP BY ROLLUP (Region, Country), ROLLUP (Country)
```

Region	Country	TotalSales
North	Canada	1800
North	USA	3500
North	Total	5300
Total	Canada	3000
Total	USA	3500
Total	Total	6500
South	Argentina	2200
South	Brazil	5500
South	Total	7700
Total	Argentina	2200
Total	Brazil	5500
Total	Total	7700

Rollup by Region and Country, with multiple rollups on Country(Total As Null)

```
1 SELECT Region, Country, SUM(SalesAmount) AS TotalSales
2 FROM Sales
3 GROUP BY CUBE (Region, Country)
```

$$\{(Region, Country), (Region), (Country), ()\}$$

```
1 SELECT Region, Country, SUM(SalesAmount) AS TotalSales
2 FROM Sales
3 GROUP BY GROUPING SETS ((Region, Country))
```

$$\{(Region, Country)\}$$

Thanks
End of Chapter 5